BBM371- Data Management

Lecture 5: Index Files and Comparing Cost Models

Indexes

- An *index* on a file speeds up selections on the *search key fields* for the index.
  - ►Any subset of the fields of a relation can be the search key for an index on the relation.
  - ►*Search key* is not the same as *key* (minimal set of fields that uniquely identify a record in a relation).
- An index contains a collection of *data entries*, and supports efficient retrieval of all data entries **k*** with a given key value **k**.
  - ► Given data entry k*, we can find record with key k in at most one disk I/O. (Details soon …)

Alternatives for Data Entry k* in Index

- In a data entry k* we can store:
  - Data record with key value **k,** or
  - <**k**, rid of data record with search key value **k**>, or
  - <**k**, list of rids of data records with search key **k**>

- Choice of alternative for data entries is orthogonal to the indexing technique used to locate data entries with a given key value **k**.
  - Examples of indexing techniques: B+ trees, hash-based structures
  - Typically, index contains auxiliary information that directs searches to the desired data entries

Alternatives for Data Entries (Contd.)

► Alternative 1:
  - ► If this is used, index structure is a file organization for data records (instead of a Heap file or sorted file).
  - ► At most one index on a given collection of data records can use Alternative 1. (Otherwise, data records are duplicated, leading to redundant storage and potential inconsistency.)
  - ► If data records are very large, # of pages containing data entries is high. Implies size of auxiliary information in the index is also large, typically.

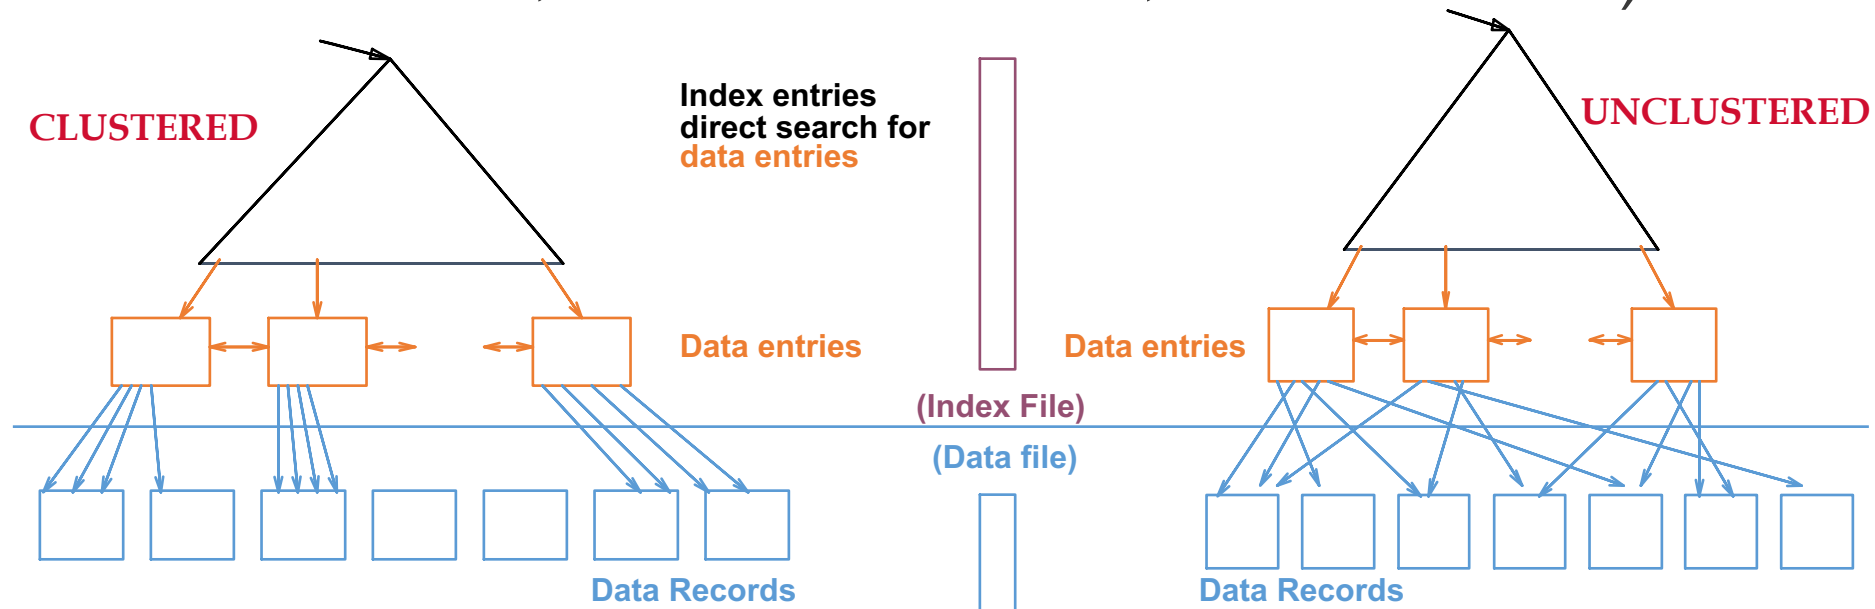# Alternatives for Data Entries (Contd.)

► Alternatives 2 and 3:

  ► Data entries typically much smaller than data records. So, better than Alternative 1 with large data records, especially if search keys are small. (Portion of index structure used to direct search, which depends on size of data entries, is much smaller than with Alternative 1.)

  ► Alternative 3 more compact than Alternative 2, but leads to variable sized data entries even if search keys are of fixed length.

Index Classification

- *Primary vs. secondary:* If search key contains primary key, then called primary index.
  - *Unique* index: Search key contains a candidate key.
- *Clustered vs. unclustered:* If order of data records is the same as, or `close to', order of data entries, then called clustered index.
  - Alternative 1 implies clustered; in practice, clustered also implies Alternative 1 (since sorted files are rare).
  - A file can be clustered on at most one search key.
  - Cost of retrieving data records through index varies *greatly* based on whether index is clustered or not!

Clustered vs. Unclustered Index

▶ Suppose that Alternative (2) is used for data entries, and that the data records are stored in a Heap file.

- ▶ To build clustered index, first sort the Heap file (with some free space on each page for future inserts).
- ▶ Overflow pages may be needed for inserts. (Thus, order of data recs is `close to', but not identical to, the sort order.)

CLUSTERED

Index entries direct search for data entries

Data entries

UNCLUSTERED

Data entries

(Index File)

(Data file)

Data Records

Data Records

Comparing Storage Techniques

Cost Model for Our Analysis

We ignore CPU costs, for simplicity:

- ►**B:** The number of data pages
- ►**R:** Number of records per page
- ►**D:** (Average) time to read or write disk page
- ►Measuring number of page I/O's ignores gains of pre-fetching a sequence of pages; thus, even I/O cost is only approximated.
- ►Average-case analysis; based on several simplistic assumptions.

☞ *Good enough to show the overall trends!*

Operations to Compare

- ▶ Scan: Fetch all records from disk

- ▶ Equality search

- ▶ Range selection

- ▶ Insert a record

- ▶ Delete a record

Assumptions in Our Analysis

▸ Single record insert and delete.

▸ Heap Files:

  ▸ Equality selection on key; exactly one match.

  ▸ Insert always at end of file.

▸ Sorted Files:

  ▸ Files compacted after deletions.

  ▸ Selections on sort field(s).

▸ Hashed Files:

  ▸ No overflow buckets, 80% page occupancy.

Cost of Operations

	Heap File	Sorted File	Hashed File
Scan all recs	**BD**	**BD**	**1.25 BD**
Equality Search	**0.5 BD**	**D $\log_2 B$**	**D**
Range Search	**BD**	**D ($\log_2 B$ + # of pages with matches)**	**1.25 BD**
Insert	**2D**	**Search + BD**	**2D**
Delete	**Search + D**	**Search + BD**	**2D**

Summary

▶ Many alternative file organizations exist, each appropriate in some situation.

▶ If selection queries are frequent, sorting the file or building an *index* is important.

▶ Index is a collection of data entries plus a way to quickly find entries with given key values.

▶ Data entries can be actual data records, <key, rid> pairs, or <key, rid-list> pairs.

  ▶ Choice orthogonal to *indexing technique* used to locate data entries with a given key value.

Summary (Contd.)

▸ Can have several indexes on a given file of data records, each with a different search key.

▸ Indexes can be classified as clustered vs. unclustered, primary vs. secondary, and dense vs. sparse. Differences have important consequences for utility/performance.

▸ Understanding the nature of the *workload* for the application, and the performance goals, is essential to developing a good design.

  ▶ What are the important queries and updates? What attributes/relations are involved?

Summary (Contd.)

- Indexes must be chosen to speed up important queries (and perhaps some updates!).
  - ▶ Index maintenance overhead on updates to key fields.
  - ▶ Choose indexes that can help many queries, if possible.
  - ▶ Build indexes to support index-only strategies.
  - ▶ Clustering is an important decision; only one index on a given relation can be clustered!
  - ▶ Order of fields in composite index key can be important.