


# BBM406: Fundamentals of Machine Learning

## Support Vector Machines


# Bias vs. Variance Trade off




- High bias (and low variance) indicates underfitting problem
- High variance (and low bias) indicates underfitting problem

# Bias vs. Variance Trade off

K=1



K=15



Figures from Hastie, Tibshirani and Friedman (Elements of Statistical Learning)

Overfitting  
(High Variance)

Underfitting  
(High bias)

# Regularization

- Linear Regression Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


- Regularized Linear Regression Cost Function

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^d \theta_j^2 \quad \left. \vphantom{\sum_{j=1}^d} \right\} \text{Regularization Parameter}$$


# Strengths of SVMs

- Good generalization
  - in theory
  - in practice
- Works well with few training instances
- Find globally best model
- Efficient algorithms
- Amenable to the kernel trick


# Intuitions



# Intuitions




# Intuitions







# Intuitions




# A “Good” Separator




# Noise in the Observations




# Ruling Out Some Separators




# Lots of Noise



# Only One Separator Remains



# Maximizing the Margin



# Why Maximize Margin

Increasing margin reduces *capacity*

- i.e., fewer possible models

Lesson from Learning Theory:

- If the following holds:
  - $H$  is sufficiently constrained in size
  - and/or the size of the training data set  $n$  is large,then low training error is likely to be evidence of low generalization error




# Alternate View of Logistic Regression

Cost of a sample:  $y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$


$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$z = \theta^T x$$

If  $y = 1$  (we want  $\theta^T x \gg 0$ )



If  $y = 0$  (we want  $\theta^T x \ll 0$ )



# Logistic Regression to SVMs

Logistic Regression:

$$\min_{\theta} - \sum_{i=1}^n (y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))) + \frac{\lambda}{2} \sum_{j=1}^d \theta_j^2$$

Support Vector Machines:

$$\min_{\theta} C \sum_{i=1}^n (y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)})) + \frac{1}{2} \sum_{j=1}^d \theta_j^2$$


You can think of  $C$  as similar to  $\frac{1}{\lambda}$

# From Logistic Regression to SVM


## Support Vector Machines:

$$\min_{\theta} C \sum_{i=1}^n (y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)})) + \frac{1}{2} \sum_{j=1}^d \theta_j^2$$

If  $y = 1$  (we want  $\theta^T x \gg 0$ )




If  $y = 0$  (we want  $\theta^T x \ll 0$ )




# Support Vector Machine

$$\min_{\theta} C \sum_{i=1}^n (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) + \frac{1}{2} \sum_{j=1}^d \theta_j^2$$

If  $y = 1$  (we want  $\theta^T x \geq 1$ )




If  $y = 0$  (we want  $\theta^T x \leq -1$ )



# Support Vector Machine


$$\min_{\theta} C \sum_{i=1}^n \underbrace{(y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)}))}_0 + \frac{1}{2} \sum_{j=1}^d \theta_j^2$$




$$\min_{\theta} \frac{1}{2} \sum_{j=1}^d \theta_j^2$$

$$\text{s.t. } \begin{aligned} \theta^T x^{(i)} &\geq 1 && \text{if } y^{(i)} = 1 \\ \theta^T x^{(i)} &\leq -1 && \text{if } y^{(i)} = 0 \end{aligned}$$


# Maximum Margin Hyperplane



# Support Vectors




# Large Margin Classifier in Presence of Outliers



$$\min_{\theta} C \sum_{i=1}^n (y^{(i)} cost_1(\theta^T x^{(i)}) + (1 - y^{(i)}) cost_0(\theta^T x^{(i)})) + \frac{1}{2} \sum_{j=1}^d \theta_j^2$$



# Vector Inner Product



$$\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \quad \mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$


$$\begin{aligned} \|\mathbf{u}\|_2 &= \sqrt{u_1^2 + u_2^2} \\ &= \text{length}(\mathbf{u}) \end{aligned}$$

$$\begin{aligned} \mathbf{u}^T \mathbf{v} &= \mathbf{v}^T \mathbf{u} \\ &= u_1 v_1 + u_2 v_2 \\ &= \|\mathbf{u}\|_2 \|\mathbf{v}\|_2 \cos \theta \\ &= p \|\mathbf{u}\|_2 \text{ where } p = \|\mathbf{v}\|_2 \cos \theta \end{aligned}$$

# Understanding the Hyperplane

$$\begin{aligned} \min_{\theta} \quad & \frac{1}{2} \sum_{j=1}^d \theta_j^2 \\ \text{s.t.} \quad & \theta^T x^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1 \\ & \theta^T x^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0 \end{aligned}$$

Assume  $\theta_0 = 0$  so that the hyperplane is centered at the origin, and that  $d = 2$



$$\begin{aligned} \theta^T x &= \|\theta\|_2 \|x\|_2 \cos \theta \\ &= p \|\theta\|_2 \end{aligned}$$

# Maximizing the Margin


$$\min_{\theta} \frac{1}{2} \sum_{j=1}^d \theta_j^2$$

$$\text{s.t. } \theta^T x^{(i)} \geq 1 \quad \text{if } y^{(i)} = 1$$


$$\theta^T x^{(i)} \leq -1 \quad \text{if } y^{(i)} = 0$$

Assume  $\theta_0 = 0$  so that the hyperplane is centered at the origin, and that  $d = 2$

Let  $p_i$  be the projection of  $\mathbf{x}_i$  onto the vector  $\theta$



Since  $p$  is small, therefore  $\|\theta\|_2$  must be large to have  $p \|\theta\|_2 \geq 1$  (or  $\leq -1$ )




Since  $p$  is large,  $\|\theta\|_2$  can be smaller to have  $p \|\theta\|_2 \geq 1$  (or  $\leq -1$ )

# Size of the Margin

For the support vectors, we have  $p \|\theta\|_2 = \pm 1$

- $p$  is the length of the projection of the SVs onto  $\theta$



Therefore,

$$p = \frac{1}{\|\theta\|_2}$$

$$\text{margin} = 2p = \frac{2}{\|\theta\|_2}$$

# What if Surface is Non-Linear?






Image from <http://www.atrandomresearch.com/iclass/>


# Kernel Methods

Making the Non-Linear Linear

# When Linear Separators Fail



# Mapping into a New Feature Space



$$\varphi: X \rightarrow \hat{X} = \varphi(x)$$

- For example, with  $x^{(i)} \in R^2$

$$\varphi \left( \begin{bmatrix} x_1^{(i)} \\ x_2^{(i)} \end{bmatrix} \right) = [x_1^{(i)}, x_2^{(i)}, x_1^{(i)} \cdot x_2^{(i)}, (x_1^{(i)})^2, (x_2^{(i)})^2]$$

- Rather than run SVM on  $\mathbf{x}_i$ , run it on  $\varphi(x^{(i)})$ 
  - Find non-linear separator in input space
- What if  $\varphi(x^{(i)})$  is really big?
- Use kernels to compute it implicitly!



# Kernels

- Find kernel  $K$  such that

$$K(x^{(i)}, x^{(j)}) = \langle \varphi(x^{(i)}), \varphi(x^{(j)}) \rangle$$

- Computing  $K(x^{(i)}, x^{(j)})$  should be efficient, much more so than computing  $\varphi(x^{(i)})$  and  $\varphi(x^{(j)})$
- Use  $K(x^{(i)}, x^{(j)})$  in SVM algorithm rather than  $x^{(i)}, x^{(j)}$
- Remarkably, this is possible!

# The Kernel Trick

“Given an algorithm which is formulated in terms of a positive definite kernel  $K_1$ , one can construct an alternative algorithm by replacing  $K_1$  with another positive definite kernel  $K_2$ ”

➤ SVMs can use the kernel trick


# The Gaussian Kernel

- Also called Radial Basis Function (RBF) kernel

$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

- Has value 1 when  $\mathbf{x}_i = \mathbf{x}_j$
- Value falls off to 0 with increasing distance
- Note: Need to do feature scaling before using Gaussian Kernel

# Gaussian Kernel Example




$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

Imagine we've learned that:

$$\theta = [-0.5, 1, 1, 0]$$

Predict +1 if  $\theta_0 + \theta_1 K(x, l_1) + \theta_2 K(x, l_2) + \theta_3 K(x, l_3) \geq 0$

# Gaussian Kernel Example



$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

Imagine we've learned that:


$$\theta = [-0.5, 1, 1, 0]$$

Predict +1 if  $\theta_0 + \theta_1 K(x, l_1) + \theta_2 K(x, l_2) + \theta_3 K(x, l_3) \geq 0$

- For  $x^{(1)}$ , we have  $K(x^{(1)}, l_1) \approx 1$ , other similarities  $\approx 0$

$$\begin{aligned}\theta_0 + \theta_1 \cdot 1 + \theta_2 \cdot 0 + \theta_3 \cdot 0 &= -0.5 + 1.1 + 1.0 + 0.1 \\ &= 0.5 \geq 0, \text{ so predict +1}\end{aligned}$$

# Gaussian Kernel Example



$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

Imagine we've learned that:


$$\theta = [-0.5, 1, 1, 0]$$

Predict +1 if  $\theta_0 + \theta_1 K(x, l_1) + \theta_2 K(x, l_2) + \theta_3 K(x, l_3) \geq 0$

- For  $x^{(2)}$ , we have  $K(x^{(2)}, l_3) \approx 1$ , other similarities  $\approx 0$

$$\begin{aligned}\theta_0 + \theta_1 \cdot 0 + \theta_2 \cdot 0 + \theta_3 \cdot 1 &= -0.5 + 1.0 + 1.0 + 0.1 \\ &= -0.5 \leq 0, \text{ so predict } -1\end{aligned}$$

# Gaussian Kernel Example



$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

Imagine we've learned that:

$$\theta = [-0.5, 1, 1, 0]$$


Predict +1 if  $\theta_0 + \theta_1 K(x, l_1) + \theta_2 K(x, l_2) + \theta_3 K(x, l_3) \geq 0$

Rough sketch of decision surface



# The Gaussian Kernel

- Also called Radial Basis Function (RBF) kernel

$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$



lower bias,  
higher variance



higher bias,  
lower variance





# Other Kernels

- Sigmoid Kernel

$$K(x^{(i)}, x^{(j)}) = \tanh(\alpha (x^{(i)})^T x^{(j)} + c)$$

- Neural networks use sigmoid as activation function
- SVM with a sigmoid kernel is equivalent to 2-layer perceptron

- Cosine Similarity Kernel

$$K(x^{(i)}, x^{(j)}) = \frac{(x^{(i)})^T x^{(j)}}{\|x^{(i)}\| \cdot \|x^{(j)}\|}$$

- Popular choice for measuring similarity of text documents
- $L_2$  norm projects vectors onto the unit sphere; their dot product is the cosine of the angle between the vectors

# Other Kernels


- Chi-squared Kernel

$$K(x^{(i)}, x^{(j)}) = \exp \left( -\gamma \sum_k \frac{(x_k^{(i)} - x_k^{(j)})^2}{x_k^{(i)} + x_k^{(j)}} \right)$$

- Widely used in computer vision applications
- Chi-squared measures distance between probability distributions
- Data is assumed to be non-negative, often with  $L_1$  norm of 1

- String kernels
- Tree kernels
- Graph kernels

# Multi-Class Classification with SVMs



$$y \in \{1, \dots, K\}$$

- Many SVM packages already have multi-class classification built in
- Otherwise, use one-vs-rest
  - Train  $K$  SVMs, each picks out one class from rest, yielding  $\theta^{(1)}, \dots, \theta^{(K)}$
  - Predict class  $i$  with largest  $(\theta^{(i)})^T x$

# Practical Advice for Applying SVMs

- Use SVM software package to solve for parameters
  - e.g., SVMlight, libsvm, cvx (fast!), etc.
- Need to specify:
  - Choice of parameter  $C$
  - Choice of kernel function
    - Associated kernel parameters

e.g.,

$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|_2^2}{2\sigma^2}\right)$$

# SVMs vs Logistic Regression

## (Advice from Andrew Ng)

$n$  = # training examples       $d$  = # features

If  $d$  is large (relative to  $n$ ) (e.g.,  $d > n$  with  $d = 10,000$ ,  $n = 10-1,000$ )

- Use logistic regression or SVM with a linear kernel

If  $d$  is small (up to 1,000),  $n$  is intermediate (up to 10,000)

- Use SVM with Gaussian kernel

If  $d$  is small (up to 1,000),  $n$  is large (50,000+)

- Create/add more features, then use logistic regression or SVM without a kernel

Neural networks likely to work well for most of these settings, but may be slower to train

# Other SVM Variations

- nu SVM
  - nu parameter controls:
    - Fraction of support vectors (lower bound) and misclassification rate (upper bound)
    - E.g.,  $\nu = 0.05$  guarantees that  $\geq 5\%$  of training points are SVs and training error rate is  $\leq 5\%$
  - Harder to optimize than C-SVM and not as scalable
- SVMs for regression
- One-class SVMs
- SVMs for clustering
- ...

# Conclusion

- SVMs find optimal linear separator
- The kernel trick makes SVMs learn non--linear decision surfaces
- Strength of SVMs:
  - Good theoretical and empirical performance
  - Supports many types of kernels
- Disadvantages of SVMs:
  - “Slow” to train/predict for huge data sets (but relatively fast!)
  - Need to choose the kernel (and tune its parameters)