



BBM371- Data Management

Lecture II External Sorting (multi-disk, multi-core)

27.12.2018

The Memory Hierarchy

- ▶ **Registers**, work at CPU-speed, a few hundreds of Bytes
- ▶ Different levels of **cache memory**, operating at 2 to 50 times CPU-speed, and ranging in size to a few Mbytes
- ▶ **Main memory**, operating at a few hundred times CPU-speed and comprising Gigabytes
- ▶ **Hard disk or solid state disk**, where access time is millions of cycles and size is several TBytes.
 - ▶ *The economical way to use disks is to transport data in large chunks*
- ▶ An analogous statement is true for any two adjacent levels of the hierarchy. The chunk size should be chosen such that the time for transferring a chunk is approximately equal to the time accessing a chunk.

The Parallel Disk Model of Aggarwal/Vitter

- ▶ It is usually phrased in terms of disks, but applies to any two adjacent levels of the memory hierarchy.
 - ▶ The machine has a CPU and a main memory of size M .
 - ▶ Data between main memory and disks is transferred in blocks of size B .
 - ▶ The machine has D disks that can be used in parallel.
- ▶ In one I/O-operation, one block of size B can be transferred between main memory and each disk.
- ▶ Algorithms are analyzed in terms of number of I/O-operations.

STXXL: Standard Template Library for Extra Large Data Sets

- ▶ The core of STXXL is an implementation of the C++ standard template library STL for external memory (out-of-core) computations
 - ▶ support of parallel disks
 - ▶ benefit from overlapping of I/O and computation
 - ▶ the I/O complexity of the algorithms remains optimal in most of the cases
 - ▶ asynchronous execution of the algorithmic components, enabling high-level task parallelism
- ▶ For internal computation, parallel algorithms from the MCSTL or the libstdc++ parallel mode are optionally utilized, making the algorithms inherently benefit from multi-core parallelism.

Merge Sort

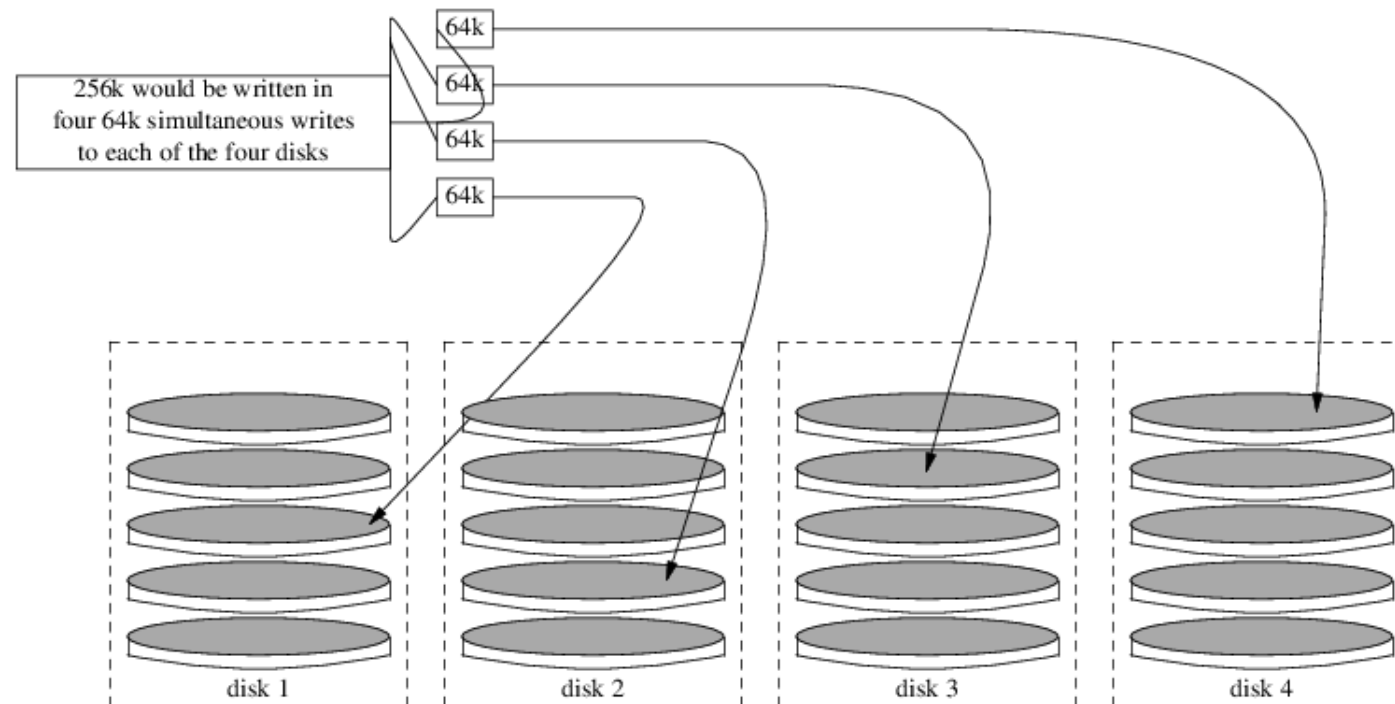
- ▶ For N items to be sorted on a buffer of size B each:
 - ▶ Merge Sort on a single disk: $2 \frac{N}{B} \left(1 + \left\lceil \log_{M/B} \frac{N}{M} \right\rceil \right) = 2 \frac{N}{B} \left\lceil \log_{M/B} \frac{N}{B} \right\rceil$ disk accesses
 - ▶ Merge Sort on D parallel disks: $2 \frac{N}{DB} \left\lceil \log_{M/(DB)} \frac{N}{DB} \right\rceil$ disk accesses
 - ▶ This looks like the internal sorting.
 - ▶ Number of blocks is $n=N/B$. Instead of binary log, we have the logarithm to the memory size measured in number of blocks.
- ▶ How good is this bound?
 - ▶ Merge Sort is optimal for one disk, but suboptimal for many disks.

Disk Striping

- ▶ We treat the D disks as a single disk with block size DB . A super-block of size DB consists of D blocks of size B .
- ▶ When a super-block is to be transferred, we transfer one standard block to each disk.
- ▶ We can generalize all single-disk results to D disks.
 - ▶ There might be more effective ways of using the D disks and that main memory can only hold $M/(DB)$ super-blocks.

Disk Striping

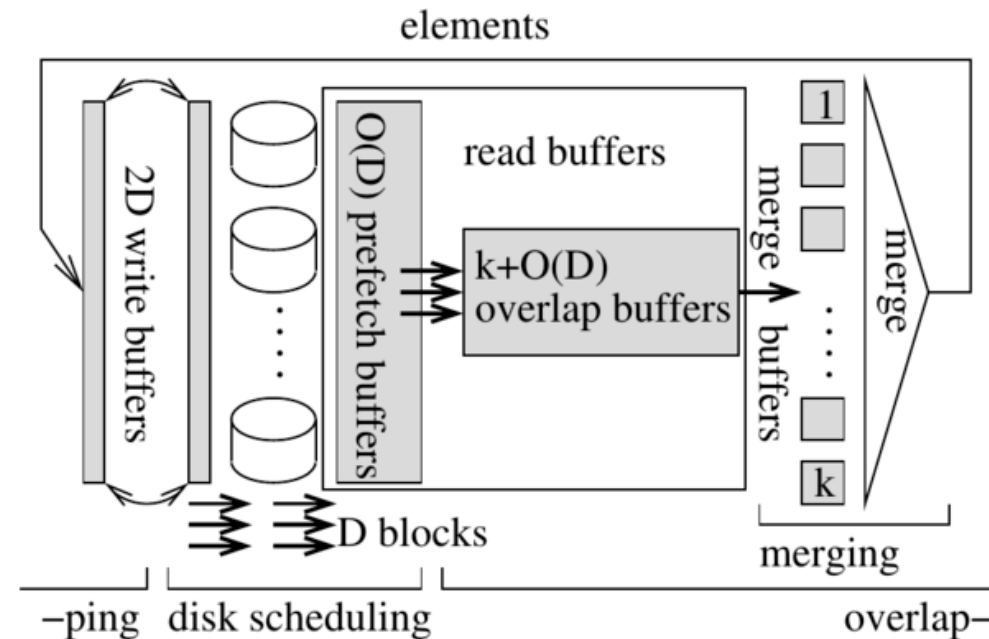
- ▶ Treat D disks as a single disk with block size DB .
- ▶ A super-block of size DB consists of D blocks of size B .



Parallel Disk Sorting

► Parallel Disk Sorting

- has an optimal I/O volume $\mathcal{O}(\frac{N}{DB} \log_{M/B} \frac{N}{B})$ (that matches the lower bound), and guarantees almost perfect overlap between I/O and computation.
- https://stxxl.org/tags/1.4.1/design_algo_sorting.html



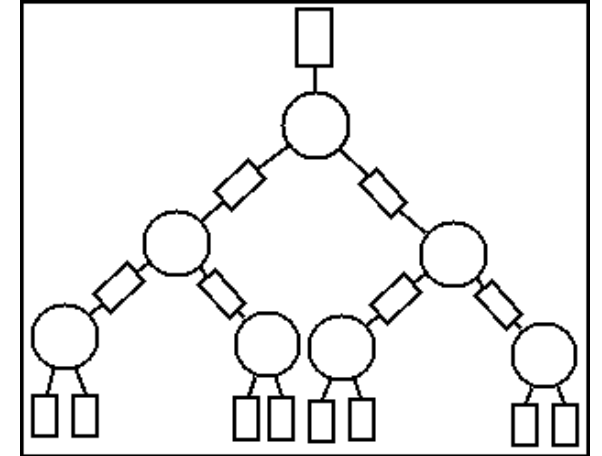
Cache-Oblivious Algorithms

▶ *Tall Cache Assumption*

- ▶ Main memory consists of M/B blocks of size B .
 - ▶ In the case of caches, M/B is called the height of the cache and B is called the width of the cache.
 - ▶ The tall cache assumption states that the height is larger than the width, i.e., $M/B \geq B$. In other words, $M \geq B^2$.
 - ▶ Many results about cache-oblivious algorithms hold true under the weaker assumption that $M \geq B^{1+\gamma}$ for some constant $\gamma > 0$.
-
- ▶ Cache-oblivious algorithms cannot use M and B in the program code. Nevertheless, they work well under the tall cache assumption.

Funnel Sort

- ▶ Funnel sort is a variant of merge sort.
- ▶ Split the input into smaller groups ●
 - ▶ Split N elements into $N^{(1/d)}$ groups of size $N^{(1-1/d)}$
 - ▶ We sort each part recursively.
 - ▶ We merge the sorted using *funnel merge* (*k-way merger*)
 - ▶ The memory layout is as in van-Emde-Boas trees..
- ▶ While being cache oblivious
 - ▶ $\Theta(n \log n)$ work
 - ▶ $\Theta((n \log n) / B)$ I/Os



Multi-Core Algorithms

- ▶ A multi-core is a parallel machine on a single chip.
 - ▶ There are several cores (= CPUs) on a single chip; up to 16 or 32 in commercial machines and up to 100 in experimental machines.
 - ▶ Each core has its own cache.
 - ▶ They share main memory
- ▶ Parallel External Memory (PEM)
 - ▶ We have P CPUs each with a private (fast) cache of size M .
 - ▶ The processors share a main memory; the main memory is unbounded in size and much slower than the private cache memories.
 - ▶ The private caches are partitioned into M/B blocks of size B each. Data is transferred in blocks between private caches and shared memory
 - ▶ In an I/O-step P blocks, one for each processor, can be transferred between main memory and private caches.
 - ▶ Concurrent read is supported. Concurrent write may or may not be supported.

Sorting

- ▶ We want to sort N elements on a PEM with P processors, each having a cache of size M . The block size is B .
 - ▶ The lower bound argument for sorting in external memory still works

$$\Omega\left(\frac{N}{PB} \log_{M/B} N/B\right)$$