

Optimization

BBM471 Database Management Systems

Dr. Fuat Akal

akal@hacettepe.edu.tr



Today's Lecture

1. Logical Optimization
2. Physical Optimization

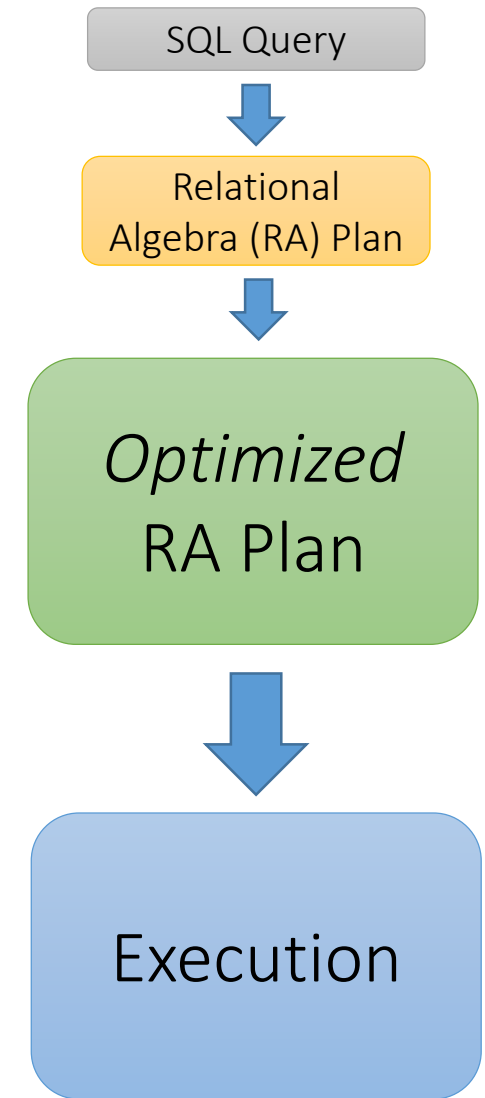
Logical vs. Physical Optimization

- **Logical optimization:**

- Find equivalent plans that are more efficient
- *Intuition: Minimize # of tuples at each step by changing the order of RA operators*

- **Physical optimization:**

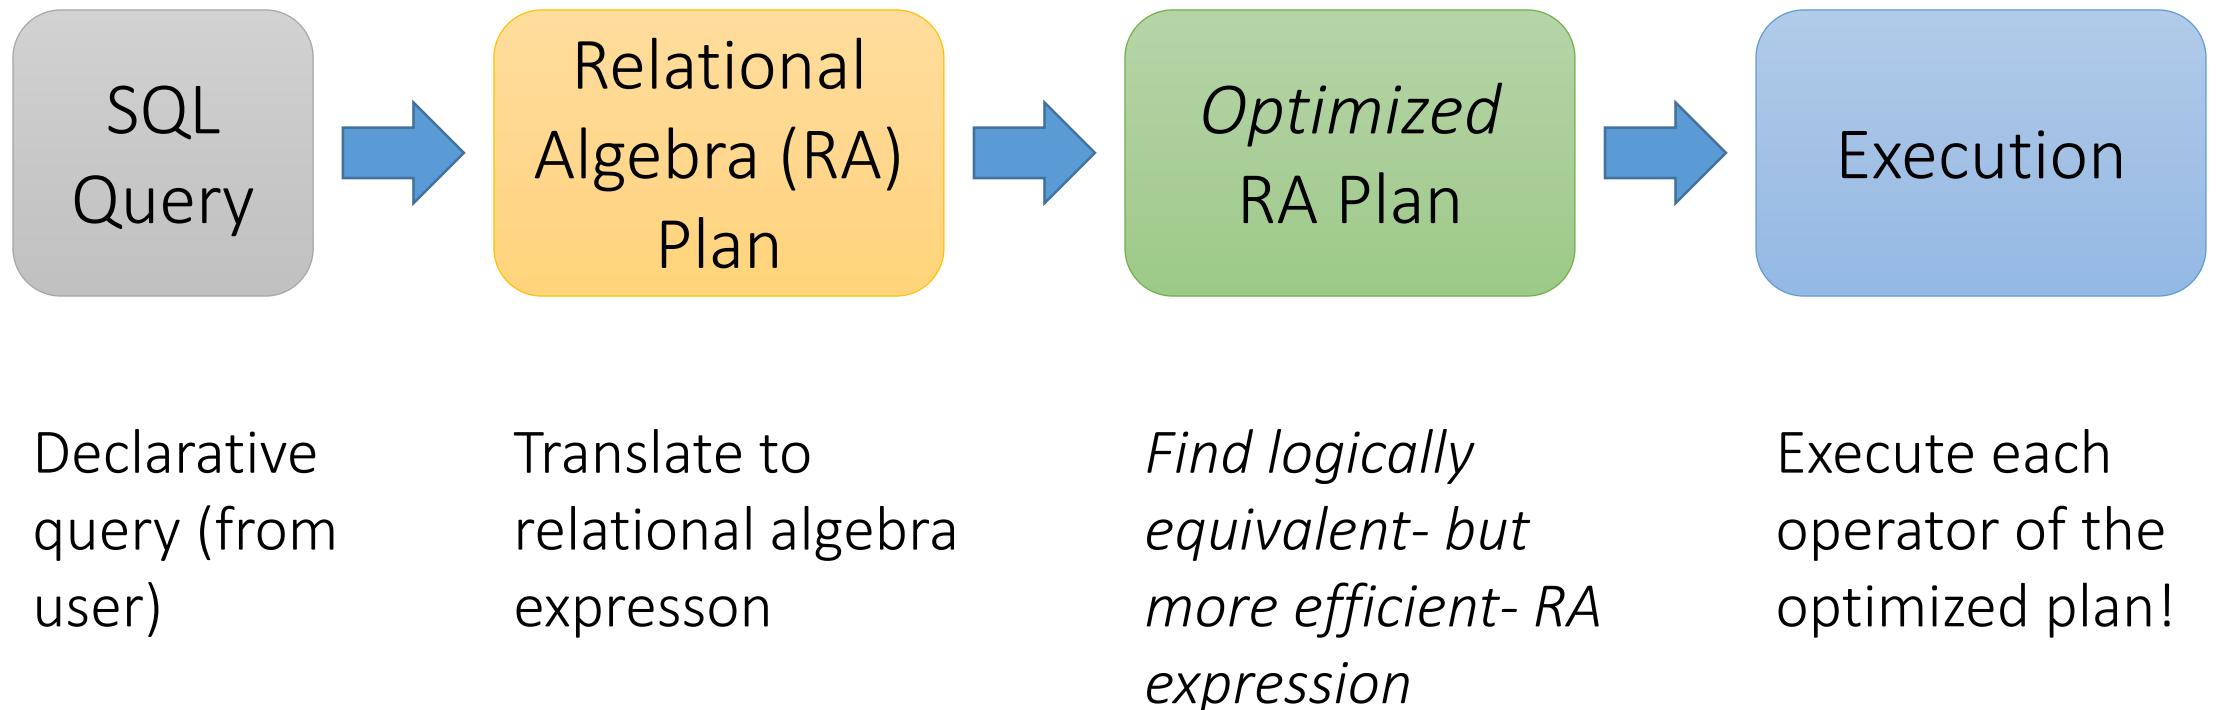
- Find algorithm with lowest IO cost to execute our plan
- *Intuition: Calculate based on physical parameters (buffer size, etc.) and estimates of data size (histograms)*



1. Logical Optimization

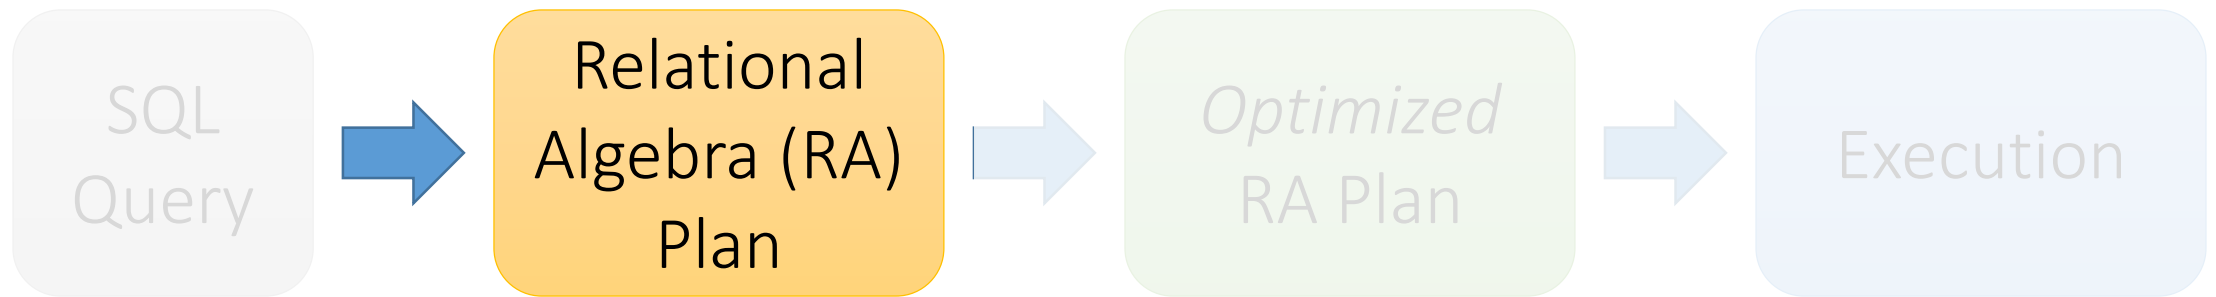
RDBMS Architecture

How does a SQL engine work ?



RDBMS Architecture

How does a SQL engine work ?



Relational Algebra allows us to translate declarative (SQL) queries into precise and optimizable expressions!

What does DBMS Do When You submit a Query?

- Translates SQL into get/put req. to backend storage
- Extracts, processes, transforms tuples from blocks
- Performs tons of optimizations
 - Choosing algorithms for SQL operators (hashing, sorting)
 - Ordering of operators (small intermediate results)
 - Semantic rewritings of queries
 - Parallel execution and concurrency
 - Load and admission control
 - Layout of data on backend storage
 - Buffer management and caching
 - ...

Input: SQL statement
Output: {tuples}

DBMS vs. OS Optimizations

- Many DBMS tasks are also carried out by OS
 - Load control
 - Buffer management
 - Access to external storage
 - Scheduling of processes
- What is the difference?
 - DBMS has intimate knowledge of workload
 - DBMS can predict and shape access pattern of a query
 - DBMS knows the contention between queries
 - OS does generic optimizations

Recall: Relational Algebra (RA)

- Five **basic** operators:

1. Selection: σ
2. Projection: Π
3. Cartesian Product: \times
4. Union: \cup
5. Difference: $-$

- Derived or auxiliary operators:

- Intersection, complement
- Joins (natural, equi-join, theta join, semi-join)
- Renaming: ρ
- Division

Recall: Converting SFW Query -> RA

Students(sid,name,gpa)
People(ssn,name,address)

```
SELECT DISTINCT  
  gpa,  
  address  
FROM Students S,  
      People P  
WHERE gpa > 3.5 AND  
      s.name = p.name;
```


$$\Pi_{gpa,address}(\sigma_{gpa>3.5}(S \bowtie P))$$

How do we represent
this query in RA?

Recall: Logical Equivalence of RA Plans

- Given relations $R(A,B)$ and $S(B,C)$:

- Here, projection & selection commute:


- $\sigma_{A=5}(\Pi_A(R)) = \Pi_A(\sigma_{A=5}(R))$

- What about here?

- $\sigma_{A=5}(\Pi_B(R)) \neq \Pi_B(\sigma_{A=5}(R))$

RDBMS Architecture



How does a SQL engine work ?



We'll look at how to then optimize these plans now


Note: We can visualize the plan as a tree

$$\Pi_B(R(A, B) \bowtie S(B, C))$$



Bottom-up tree traversal = order of operation execution!


A simple plan



What SQL query does this correspond to?

Are there any logically equivalent RA expressions?

“Pushing down” projection



Why might we prefer this plan?

Takeaways

- This process is called logical optimization
- Many equivalent plans used to search for “good plans”
- Relational algebra is an important abstraction

RA commutators

- The basic commutators:
 - Push **projection** through (1) **selection**, (2) **join**
 - Push **selection** through (3) **selection**, (4) **projection**, (5) **join**
 - *Also*: Joins can be re-ordered!
- Note that this is not an exhaustive set of operations
 - This covers *local re-writes*; *global re-writes possible but much harder*

This simple set of tools allows us to greatly improve the execution time of queries by optimizing RA plans!

Optimizing the SFW RA Plan



Translating to RA

R(A,B) S(B,C) T(C,D)

SELECT R.A,S.D
FROM R,S,T
WHERE R.B = S.B
AND S.C = T.C
AND R.A < 10;



$\Pi_{A,D}(\sigma_{A < 10}(T \bowtie (R \bowtie S)))$



Logical Optimization

- Heuristically, we want selections and projections to occur as early as possible in the plan
 - Terminology: “push down **selections**” and “pushing down **projections.**”
- **Intuition:** We will have fewer tuples in a plan.
 - Could fail if the selection condition is very expensive (say runs some image processing algorithm).
 - Projection could be a waste of effort, but more rarely.


Optimizing RA Plan

R(A,B) S(B,C) T(C,D)

```
SELECT R.A,S.D  
FROM R,S,T  
WHERE R.B = S.B  
AND S.C = T.C  
AND R.A < 10;
```

Push down
selection on A so
it occurs earlier

$\Pi_{A,D}(\sigma_{A<10}(T \bowtie (R \bowtie S)))$




Optimizing RA Plan

R(A,B) S(B,C) T(C,D)

```
SELECT R.A,S.D  
FROM R,S,T  
WHERE R.B = S.B  
AND S.C = T.C  
AND R.A < 10;
```

Push down
selection on A so
it occurs earlier

$\Pi_{A,D}(T \bowtie (\sigma_{A < 10}(R) \bowtie S))$




Optimizing RA Plan

R(A,B) S(B,C) T(C,D)

```
SELECT R.A,S.D  
FROM R,S,T  
WHERE R.B = S.B  
AND S.C = T.C  
AND R.A < 10;
```

Push down
projection so it
occurs earlier

$\Pi_{A,D}(T \bowtie (\sigma_{A<10}(R) \bowtie S))$



Optimizing RA Plan


R(A,B) S(B,C) T(C,D)

SELECT R.A,S.D
FROM R,S,T
WHERE R.B = S.B
AND S.C = T.C
AND R.A < 10;

We eliminate B
earlier!

In general, when
is an attribute not
needed...?


$\Pi_{A,D} \left(T \bowtie \Pi_{A,C} (\sigma_{A < 10}(R) \bowtie S) \right)$



Query Rewrite: Unnesting of Views

- Example: Unnesting of Views


```
select A.x  
from A  
where y in  
      (select y from B)
```



```
select A.x  
from A, B  
where A.y = B.y
```

- Example: Unnesting of Views

```
select A.x  
from A  
where exists  
      (select * from B where A.y = B.y)
```




```
select A.x  
from A, B  
where A.y = B.y
```

Query Rewrite: Predicate Augmentation

- Example: Predicate Augmentation

```
select *  
from   A, B, C  
where  A.x = B.x  
       and B.x = C.x
```



```
select *  
from   A, B, C  
where  A.x = B.x  
       and B.x = C.x  
       and A.x = C.x
```

Why is that useful?

Why Predicate Augmentation?

A (odd numbers)		B (all numbers)		C (even numbers)	
...	x	...	x	...	x
...	1	...	1	...	2
...	3	...	2	...	4
...	5	...	3	...	6
...

- $\text{Cost}((A \times C) \times B) < \text{Cost}((A \times B) \times C)$
 - get second join for free
- Query Rewrite does not know that, ...
 - but it knows that it might happen and hopes for optimizer...
- Codegen gets rid of unnecessary predicates (e.g., $A.x = B.x$)

Query Optimization

- Two tasks
 - Determine order of operators
 - Determine algorithm for each operator (hashing, sorting, ...)
- Components of a query optimizer
 - Search space
 - Cost model
 - Enumeration algorithm (NP hard)
- Working principle
 - Enumerate alternative plans
 - Apply cost model to alternative plans
 - Select plan with lowest expected cost

Optimization: Does It Really Matter? - 1

- $A \bowtie B \bowtie C$
 - $\text{size}(A) = 10,000$
 - $\text{size}(B) = 100$
 - $\text{size}(C) = 1$
 - $\text{cost}(X \bowtie Y) = \text{size}(X) + \text{size}(Y)$
- $\text{cost}((A \bowtie B) \bowtie C) = 1,010,101$
 - $\text{cost}(A \bowtie B) = 10,100$
 - $\text{cost}(X \bowtie C) = 1,000,001$ with $X = A \bowtie B$
- $\text{cost}(A \bowtie (B \bowtie C)) = 10,201$
 - $\text{cost}(B \bowtie C) = 101$
 - $\text{cost}(A \bowtie X) = 10,100$ with $X = B \bowtie C$

Optimization: Does It Really Matter? - 2

- $A \bowtie B \bowtie C$
 - $\text{size}(A) = 1000$
 - $\text{size}(B) = 1$
 - $\text{size}(C) = 1$
 - $\text{cost}(X \bowtie Y) = \text{size}(X) * \text{size}(Y)$
- $\text{cost}((A \bowtie B) \bowtie C) = 2000$
 - $\text{cost}(A \bowtie B) = 1000$
 - $\text{cost}(\mathbf{X} \bowtie C) = 1000$ with $\mathbf{X} = A \bowtie B$
- $\text{cost}(A \bowtie (B \bowtie C)) = 1001$
 - $\text{cost}(B \bowtie C) = 1$
 - $\text{cost}(A \bowtie \mathbf{X}) = 1000$ with $\mathbf{X} = B \bowtie C$

2. Physical Optimization

What you will learn about in this section

1. Index Selection
2. Histograms
3. Materialized Views

Index Selection

Input:

- Schema of the database
- **Workload description:** set of (query template, frequency) pairs

Goal: Select a set of indexes that minimize execution time of the workload.

- Cost / benefit balance: Each additional index may help with some queries, but requires updating

This is an optimization problem!

Example

Workload
description:

```
SELECT pname  
FROM Product  
WHERE year = ? AND category = ?
```

Frequency
10,000,000

```
SELECT pname,  
FROM Product  
WHERE year = ? AND Category = ?  
AND manufacturer = ?
```

Frequency
10,000,000

Which indexes might we choose?

Example

Workload
description:

```
SELECT pname  
FROM Product  
WHERE year = ? AND category =?
```

Frequency
10,000,000

```
SELECT pname  
FROM Product  
WHERE year = ? AND Category =?  
AND manufacturer = ?
```

Frequency
100

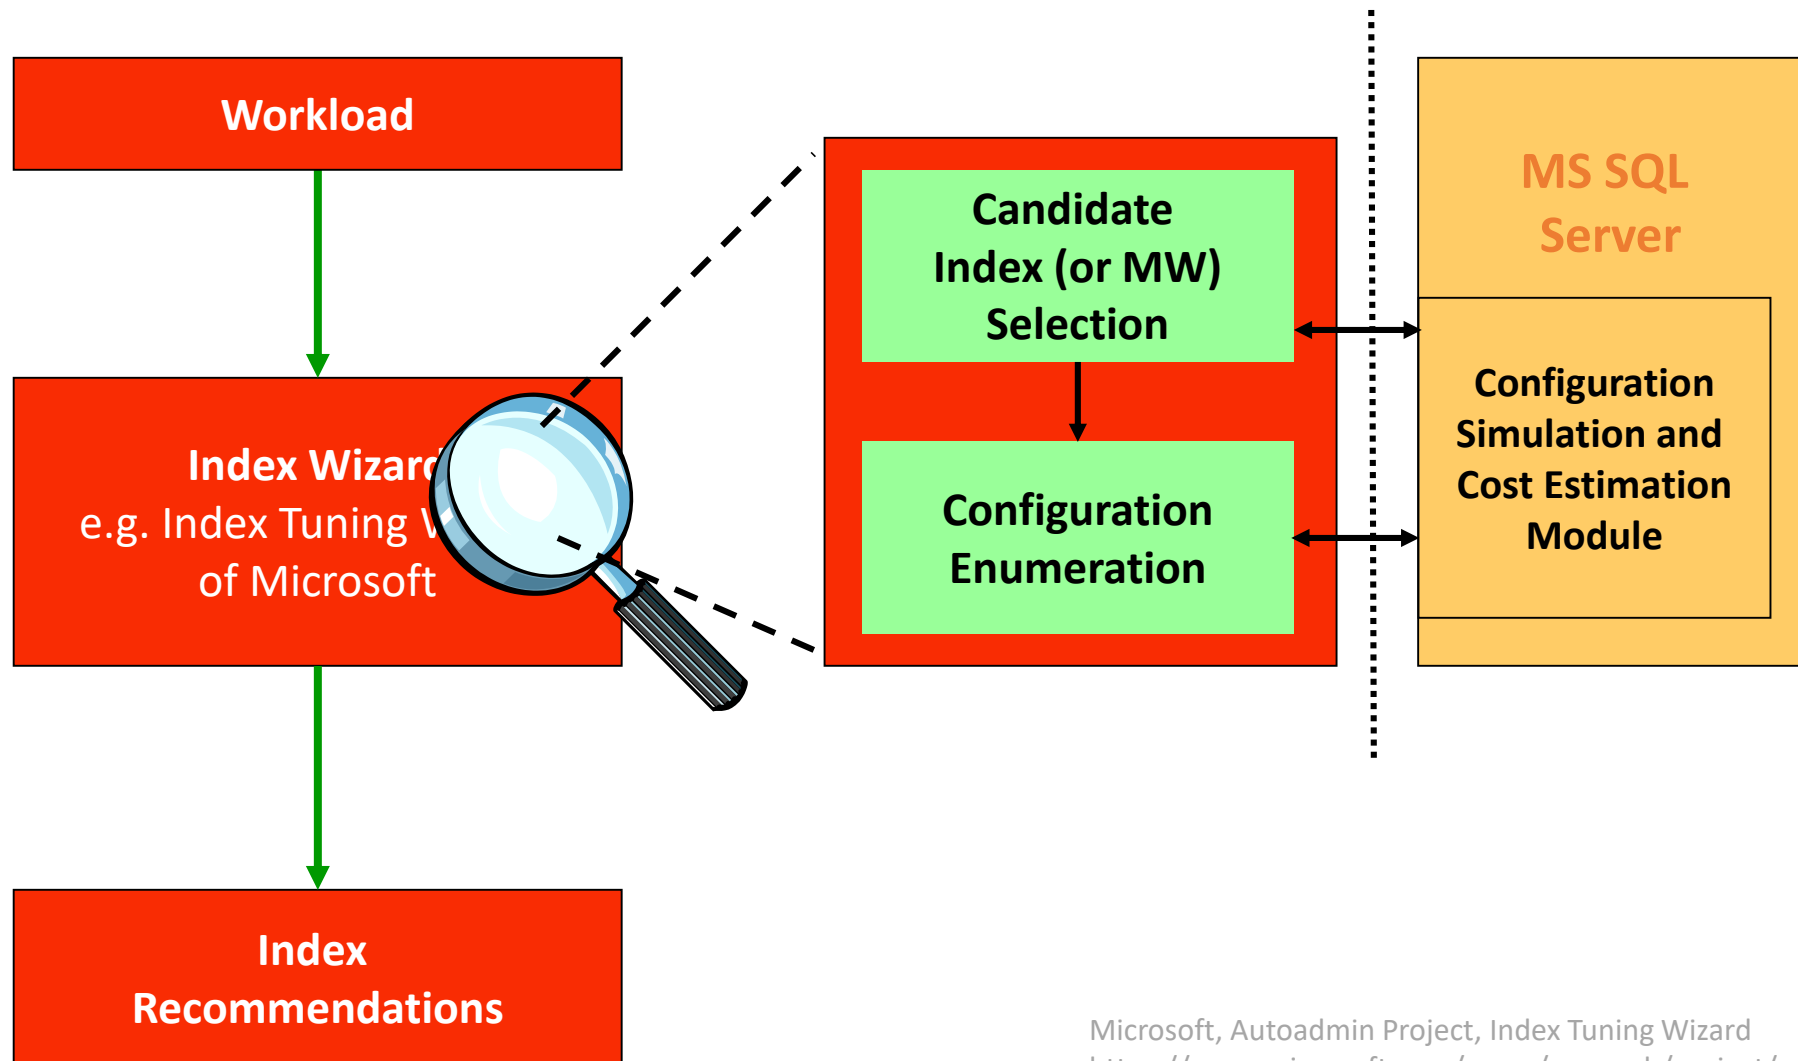
Now which indexes might we choose? Worth keeping an index with manufacturer in its search key around?

Simple Heuristic

- Can be framed as standard optimization problem: Estimate how cost changes when we add index.
 - We can ask the optimizer!
- Search over all possible space is too expensive, optimization surface is really nasty.
 - Real DBs may have 1000s of tables!
- Techniques to exploit *structure* of the space.
 - In SQL Server Autoadmin.

NP-hard problem, but can be solved!

Automatic Database Design (Physical Layout)



Microsoft, Autoadmin Project, Index Tuning Wizard
<https://www.microsoft.com/en-us/research/project/autoadmin/>

Estimating index cost?

- Note that to frame as optimization problem, we first need an estimate of the **cost** of an index lookup
- Need to be able to estimate the costs of different indexes / index types...

We will see this mainly depends on getting estimates of result set size!

Histograms & IO Cost Estimation

IO Cost Estimation via Histograms

- For **index selection**:
 - What is the cost of an index lookup?
- Also for **deciding which algorithm to use**:
 - Ex: To execute $R \bowtie S$, which join algorithm should DBMS use?
 - What if we want to compute $\sigma_{A>10}(R) \bowtie \sigma_{B=1}(S)$?
- In general, we will need some way to ***estimate intermediate result set sizes***


Histograms provide a way to efficiently store estimates of these quantities

Histograms

- A histogram is a set of value ranges (“buckets”) and the frequencies of values in those buckets occurring
- Can be used to estimate cardinality of result sets
- How to choose the buckets?
 - Equiwidth & Equidepth
- Turns out high-frequency values are **very** important

Example Histogram


Frequency



How do we
compute how
many values
between 8 and
10?
(Yes, it's obvious)

Problem: counts take up too much space!

Full vs. Uniform Counts



How much space do the full counts (bucket_size=1) take?


How much space do the uniform counts (bucket_size=ALL) take?

Fundamental Tradeoffs

- Want high resolution (like the full counts)
- Want low space (like uniform)
- Histograms are a compromise!


So how do we compute the “bucket” sizes?

Equi-width



All buckets roughly the same width

Equidepth



All buckets contain roughly the same number of items (total frequency). Able to adapt to skew.


Histograms

- Simple, intuitive and popular
- Parameters: # of buckets and type
- Can extend to many attributes (multidimensional)

Maintaining Histograms

- Histograms require that we update them!
 - Typically, you must run/schedule a command to update statistics on the database
 - Out of date histograms can be terrible!
- There is research work on self-tuning histograms and the use of query feedback
 - Oracle 11g

Nasty example



1. we insert many tuples with value > 16
2. we do **not** update the histogram
3. we ask for values > 20 ?

Precomputing Things

Views

- Views are relations, except that they are not physically stored.
- For presenting different information to different users
- **Employee** (ssn, name, department, project, salary)

```
CREATE VIEW Developers AS
  SELECT name, project
  FROM Employee
  WHERE department = "Development"
```

- Set privileges so that Payroll has access to **Employee**, others only to **Developers**

Example View Based on a Join

```
CREATE VIEW Seattle-view AS
  SELECT buyer, seller, product, store
  FROM Person, Purchase
  WHERE Person.city = "Seattle"
        AND Person.name = Purchase.buyer
```

We have a new virtual table:

Seattle-view (buyer, seller, product, store)

```
SELECT name, store
FROM Seattle-view, Product
WHERE Seattle-view.product = Product.name
      AND Product.category = "shoes"
```

View is not Really a Table!

```
SELECT    name, Seattle-view.store
FROM      Seattle-view, Product
WHERE     Seattle-view.product = Product.name AND
          Product.category = "shoes"
```



This is what happens when you query a view

```
SELECT    name, Purchase.store
FROM      Person, Purchase, Product
WHERE     Person.city = "Seattle" AND
          Person.name = Purchase.buyer AND
          Purchase.product = Product.name AND
          Product.category = "shoes"
```

Pros vs. Cons of Views

+Enforce Business Rules – Use views to define business rules, such as when an item is active, or what is meant by “popular.”

+Consistency – Simplify complicated query logic and calculations by hiding it behind the view’s definition.

+Security – Restrict access to a table, yet allow users to access non-confidential data via views.

+Simplicity – Databases with many tables possess complex relationships, which can be difficult to navigate if you aren’t comfortable using Joins.

—Performance – Each time a view is referenced, the query used to define it, is rerun.

—Modifications – Not all views support INSERT, UPDATE, or DELETE operations.

Materialized Views

- Unlike views, materialized views also store the results of the query in the database.
- Designed to improve the performance of the database by doing some intensive work in advance.
 - Can be used to pre-collect aggregate values
 - Assemble data that would come from many different tables, which would in turn require many different joins to be performed

Ordinary vs. Materialized Views

- Ordinary views
 - Virtual table
 - Named select statement
- Part of the SQL standard
- Syntax
 - `CREATE VIEW viewName AS selectStatement`
- Physical table
 - Replication of master data at a single point in time
- Not part of the SQL standard
- Syntax
 - `CREATE MATERIALIZED VIEW viewName AS selectStatement`

Why Use Materialized Views?

- Replicate data to non-master sites
 - To save network traffic when data is used in transactions
- Cache expensive queries
 - Expensive in terms of time or memory
 - Example: Sum, average or other calculations on large amounts of data

Types of Materialized Views

- Read-only
 - Insert, update or delete **NOT** allowed
- Updateable
 - Insert, update and delete on the view is allowed
 - Changes made to the view are pushed back to the master tables at refresh
- Writeable
 - Insert, update and delete on the view is allowed
 - Changes made to the view are **NOT** pushed back to the master tables at refresh

Refreshing a Materialized View

- Refresh types
 - Complete refresh
 - Recreates the materialized view
 - Fast (Incremental) refresh
 - Only changed data is refreshed
- Initiating a refresh
 - Scheduled refresh
 - On-demand refresh

Summary

- Database is doing lots of optimizations without you knowing
- Logical optimizations consider rewriting queries
- Physical optimizations uses additional structures to facilitate query evaluation

Acknowledgements

The course material used for this lecture is mostly taken and/or adopted from the course materials of the *CS145 Introduction to Databases* lecture given by *Christopher Ré* at *Stanford University* (<http://web.stanford.edu/class/cs145/>).