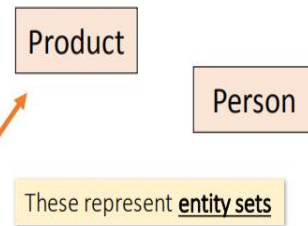


Ders02 - ER Model:

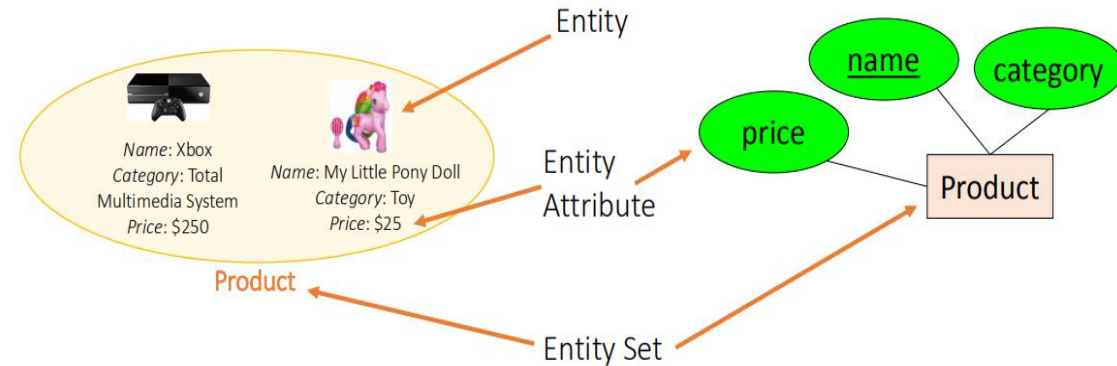
Entities and Entity Sets

- **Entities & entity sets** are the primitive unit of the E/R model
 - Entities are the individual objects, which are members of entity sets
 - Ex: A specific person or product
 - Entity sets are the *classes* or *types* of objects in our model
 - Ex: Person, Product
 - *These are what is shown in E/R diagrams - as rectangles*
 - *Entity sets represent the sets of all possible entities*



Entities vs. Entity Sets


Example:



Keys

- A key is a **minimal** set of attributes that uniquely identifies an entity.

Denote elements of the primary key by underlining.




The E/R model forces us to designate a single **primary** key, though there may be multiple candidate keys

Here, {price, category} is not a key.

If it were, what would it mean?

What is a Relationship?




A **relationship** between entity sets P and C is a *subset of all possible pairs of entities in P and C*, with tuples uniquely identified by *P and C's keys*

What is a Relationship?


- There can only be **one relationship for every unique combination of entities**
- This also means that **the relationship is uniquely determined by the keys of its entities**
- Example: the “key” for Makes (to right) is {Product.name, Company.name}

This follows from our mathematical definition of a relationship- it's a SET!




Multiplicity of E/R Relationships

One-to-one:




Indicated using arrows


Many-to-one:



One-to-many:



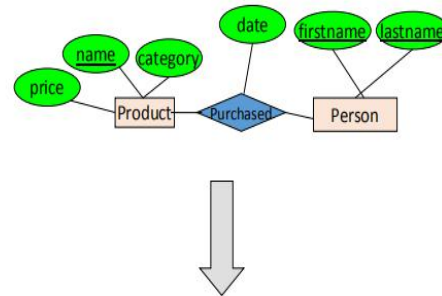
Many-to-many:



$X \rightarrow Y$ means there exists a function mapping from X to Y (recall the definition of a function)

From E/R Diagrams to Relational Schema

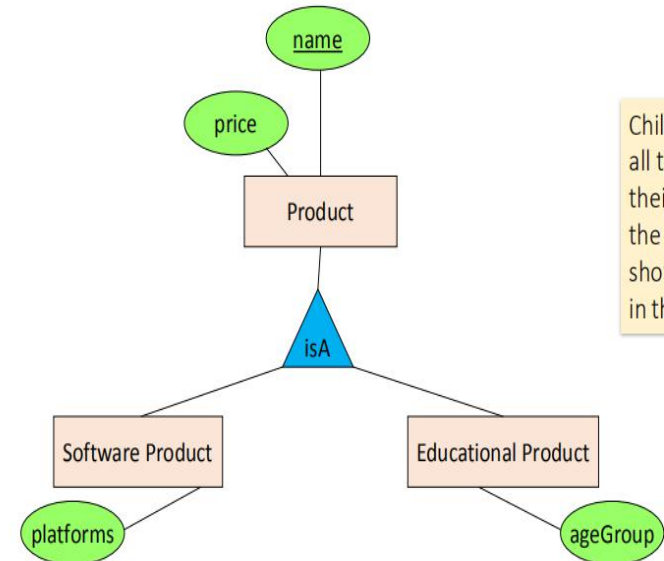
```
CREATE TABLE Purchased(
  name CHAR(50),
  firstname CHAR(50),
  lastname CHAR(50),
  date DATE,
  PRIMARY KEY (name, firstname, lastname),
  FOREIGN KEY (name)
    REFERENCES Product,
  FOREIGN KEY (firstname, lastname)
    REFERENCES Person
)
```



Purchased

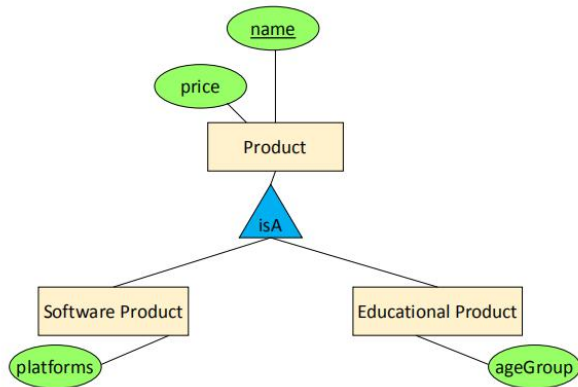
name	firstname	lastname	date
Gizmo1	Bob	Joe	01/01/15
Gizmo2	Joe	Bob	01/03/15
Gizmo1	JoeBob	Smith	01/05/15

Modeling Subclasses



Child subclasses contain all the attributes of *all* of their parent classes **plus** the new attributes shown attached to them in the E/R diagram

Think like tables...



Product

<u>name</u>	price	category
Gizmo	99	gadget
Camera	49	photo
Toy	39	gadget


Sw.Product

<u>name</u>	platforms
Gizmo	unix


Ed.Product

<u>name</u>	ageGroup
Gizmo	todler
Toy	retired

Participation Constraints: Partial v. Total



Are there products made by no company?
Companies that don't make a product?



Bold line indicates total participation (i.e. here: all products are made by a company)

Weak Entity Sets

Entity sets are weak when their key comes from other classes to which they are related.



E/R Summary

- E/R diagrams are a visual syntax that allows technical and non-technical people to talk
 - For conceptual design
- Basic constructs: **entity**, **relationship**, and **attributes**
- A good design is faithful to the constraints of the application, but not overzealous

❖ SQL is a...

- ✓ Data Definition Language (DDL)
 - Define relational schemata
 - Create/alter/delete tables and their attributes
- ✓ Data Manipulation Language (DML)
 - Insert/delete/modify tuples in tables
 - Query one or more tables – discussed next!

❖ Data Types in SQL

- Atomic types:
- ✓ Characters: CHAR(20), VARCHAR(50)
 - ✓ Numbers: INT, BIGINT, SMALLINT, FLOAT
 - ✓ Others: MONEY, DATETIME, ...

Every attribute must have an atomic type. Hence tables are flat.

❖ Table Schemas

- ✓ The schema of a table is the table name, its attributes, and their types:

Product(Pname: string, Price: float, Category: string, Manufacturer: string)

- ✓ A key is an attribute whose values are unique; we underline a key:

Product(Pname: string, Price: float, Category: string, Manufacturer: string)

SQL Query

- Basic form (there are many many more bells and whistles)

```
SELECT <attributes>
FROM <one or more relations>
WHERE <conditions>
```

Call this a SFW query.

Tables in SQL

Product

Attribute Column		
PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

Tuple || Row

A relation or table is a multiset of tuples having the attributes specified by the schema

The number of tuples is the cardinality of the relation

A multiset is an unordered list (or: a set with multiple duplicate instances allowed)

List: [1, 1, 2, 3]
Set: {1, 2, 3}
Multiset: {1, 1, 2, 3}

i.e. no next(), etc. methods!

The number of attributes is the arity of the relation

❖ A Few Details

- ✓ SQL statements are case insensitive: “Same: SELECT, Select, select” or “Same: Product, product”
- ✓ Values are not: “Different: ‘Seattle’, ‘seattle’ ”
- ✓ Use single quotes for constants: ‘abc’ - yes but “abc” - no

LIKE: Simple String Pattern Matching DISTINCT: Eliminating Duplicates

- s **LIKE** p : pattern matching on strings
- p may contain two special symbols:
 - % = any sequence of characters
 - _ = any single character

```
SELECT DISTINCT Category
FROM Product
```



Category
Gadgets
Photography
Household

Versus

```
SELECT *
FROM Products
WHERE PName LIKE '%gizmo%'
```

```
SELECT Category
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

ORDER BY: Sorting the Results

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Category='gizmo' AND Price > 50
ORDER BY Price, PName
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the **DESC** keyword.

Foreign Key constraints

- Suppose we have the following schema:

```
Students(sid: string, name: string, gpa: float)
Enrolled(student_id: string, cid: string, grade: string)
```

- And we want to impose the following constraint:

- 'Only bona fide students may enroll in courses' i.e. a student must appear in the Students table to enroll in a class

Students			Enrolled		
sid	name	gpa	student_id	cid	grade
101	Bob	3.2	123	564	A
123	Mary	3.8	123	537	A+

student_id alone is not a key- what is?

We say that student_id is a **foreign key** that refers to Students

Declaring Foreign Keys

```
Students(sid: string, name: string, gpa: float)
Enrolled(student_id: string, cid: string, grade: string)

CREATE TABLE Enrolled(
    student_id CHAR(20),
    cid CHAR(20),
    grade CHAR(10),
    PRIMARY KEY (student_id, cid),
    FOREIGN KEY (student_id) REFERENCES Students(sid)
)
```

Joins

```
Product(PName, Price, Category, Manufacturer)
Company(CName, StockPrice, Country)
```

Several equivalent ways to write a basic join in SQL:

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

```
SELECT PName, Price
FROM Product
JOIN Company ON Manufacturer = CName
              AND Country='Japan'
WHERE Price <= 200
```

Both equivalent ways to resolve variable ambiguity

```
Person(name, address, worksfor)
Company(name, address)
```

```
SELECT DISTINCT Person.name, Person.address
FROM      Person, Company
WHERE     Person.worksfor = Company.name
```

```
SELECT DISTINCT p.name, p.address
FROM      Person p, Company c
WHERE     p.worksfor = c.name
```

Tuple Variable Ambiguity in Multi-Table

An Unintuitive Query

```
SELECT DISTINCT R.A
FROM R, S, T
WHERE R.A=S.A OR R.A=T.A
```

- Recall the semantics!

1. Take cross-product
2. Apply selections / conditions
3. Apply projection

- If $S = \{\}$, then the cross product of $R, S, T = \{\}$, and the query result = $\{\}$!

Must consider semantics here.

Are there more explicit way to do set operations like this?


Explicit Set Operators: INTERSECT

```
SELECT R.A
FROM R, S
WHERE R.A=S.A
```

INTERSECT

```
SELECT R.A
FROM R, T
WHERE R.A=T.A
```

$\{r.A \mid r.A = s.A\} \cap \{r.A \mid r.A = t.A\}$




UNION

```
SELECT R.A
FROM R, S
WHERE R.A=S.A
```

UNION

```
SELECT R.A
FROM R, T
WHERE R.A=T.A
```

$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$



Why aren't there duplicates?

What if we want duplicates?


UNION ALL

```
SELECT R.A
FROM R, S
WHERE R.A=S.A
```

UNION ALL

```
SELECT R.A
FROM R, T
WHERE R.A=T.A
```

$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$



*ALL indicates
Multiset
operations*

EXCEPT

```
SELECT R.A
FROM R, S
WHERE R.A=S.A
```

EXCEPT

```
SELECT R.A
FROM R, T
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \setminus \{r.A \mid r.A = t.A\}$$


What is the multiset version?

INTERSECT: Remember the semantics!

```
Company(name, hq_city) AS C
Product(pname, maker, factory_loc) AS P
```

```
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='US'
```

```
INTERSECT
SELECT hq_city
FROM Company, Product
WHERE maker = name
AND factory_loc='China'
```

Example: C JOIN P on maker = name

C.name	C.hq_city	P.pname	P.maker	P.factory_loc
X Co.	Seattle	X	X Co.	U.S.
Y Inc.	Seattle	X	Y Inc.	China

X Co has a factory in the US (but not China)
Y Inc. has a factory in China (but not US)

But Seattle is returned by the query!

We did the INTERSECT on the wrong attributes!

One Solution: Nested Queries

```
Company(name, hq_city)
Product(pname, maker, factory_loc)
```

```
SELECT DISTINCT hq_city
FROM Company, Product
WHERE maker = name
AND name IN (
    SELECT maker
    FROM Product
    WHERE factory_loc = 'US')
AND name IN (
    SELECT maker
    FROM Product
    WHERE factory_loc = 'China')
```

"Headquarters of companies which make products in US **AND** China"

Note: If we hadn't used DISTINCT here, how many copies of each hq_city would have been returned?

Nested Queries

```
SELECT DISTINCT c.city
FROM Company c,
Product pr,
Purchase p
WHERE c.name = pr.maker
AND pr.name = p.product
AND p.buyer = 'Joe Blow'
```

```
SELECT DISTINCT c.city
FROM Company c
WHERE c.name IN (
    SELECT pr.maker
    FROM Purchase p, Product pr
    WHERE p.product = pr.name
    AND p.buyer = 'Joe Blow')
```

Now they are equivalent

Subqueries Returning Relations

You can also use operations of the form:

- $s > \text{ALL } R$
- $s < \text{ANY } R$
- $\text{EXISTS } R$

ANY and ALL not supported by SQLite.

Ex: `Product(name, price, category, maker)`

```
SELECT name
FROM Product
WHERE price > ALL(
    SELECT price
    FROM Product
    WHERE maker = 'Gizmo-Works')
```

Find products that are more expensive than all those produced by "Gizmo-Works"

Subqueries Returning Relations

You can also use operations of the form:

- $s > \text{ALL } R$
- $s < \text{ANY } R$
- $\text{EXISTS } R$

Ex: `Product(name, price, category, maker)`

```
SELECT p1.name
FROM Product p1
WHERE p1.maker = 'Gizmo-Works'
AND EXISTS(
    SELECT p2.name
    FROM Product p2
    WHERE p2.maker <> 'Gizmo-Works'
    AND p1.name = p2.name)
```

<> means !=

Find 'copycat' products, i.e. products made by competitors with the same names as products made by "Gizmo-Works"

Nested queries as alternatives to INTERSECT and EXCEPT

INTERSECT and EXCEPT not in some DBMSs!

```
(SELECT R.A, R.B
FROM R)
INTERSECT
(SELECT S.A, S.B
FROM S)
```



```
SELECT R.A, R.B
FROM R
WHERE EXISTS(
    SELECT *
    FROM S
    WHERE R.A=S.A AND R.B=S.B)
```

```
(SELECT R.A, R.B
FROM R)
EXCEPT
(SELECT S.A, S.B
FROM S)
```



```
SELECT R.A, R.B
FROM R
WHERE NOT EXISTS(
    SELECT *
    FROM S
    WHERE R.A=S.A AND R.B=S.B)
```

If R, S have no duplicates, then can write without sub-queries (HOW?)

Correlated Queries

`Movie(title, year, director, length)`

```
SELECT DISTINCT title
FROM Movie AS m
WHERE year <> ANY(
    SELECT year
    FROM Movie
    WHERE title = m.title)
```

Find movies whose title appears more than once.

Complex Correlated Query

```
Product(name, price, category, maker, year)
```

```
SELECT DISTINCT x.name, x.maker
FROM Product AS x
WHERE x.price > ALL(
    SELECT y.price
    FROM Product AS y
    WHERE x.maker = y.maker
    AND y.year < 1972)
```

Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

Can be very powerful (also much harder to optimize)

Grouping and Aggregation

```
Purchase(product, date, price, quantity)
```

```
SELECT product,
    SUM(price * quantity) AS TotalSales
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
```

Find total sales after 10/1/2005 per product.

Aggregation: COUNT

- COUNT applies to duplicates, unless otherwise stated

```
SELECT COUNT(category)
FROM Product
WHERE year > 1995
```

Note: Same as COUNT(*). Why?

We probably want:

```
SELECT COUNT(DISTINCT category)
FROM Product
WHERE year > 1995
```

HAVING Clause

```
SELECT product, SUM(price*quantity)
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
HAVING SUM(quantity) > 100
```

Same query as before, except that we consider only products that have more than 100 buyers

HAVING clauses contains conditions on **aggregates**

Whereas WHERE clauses condition on *individual tuples...*

Quantifiers

Product(name, price, company)
Company(name, city)

```
SELECT DISTINCT Company.cname
FROM Company, Product
WHERE Company.name = Product.company
AND Product.price < 100
```

An **existential quantifier** is a logical quantifier (roughly) of the form “there exists”

Find all companies that make some products with price < 100

Quantifiers

Product(name, price, company)
Company(name, city)

```
SELECT DISTINCT Company.cname
FROM Company
WHERE Company.name NOT IN(
  SELECT Product.company
  FROM Product.price >= 100)
```

A **universal quantifier** is of the form “for all”

Find all companies with products all having price < 100

↓ Equivalent

Find all companies that make only products with price < 100

Null Values

- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *
FROM Person
WHERE (age < 25)
AND (height > 6 AND weight > 190)
```

Won't return e.g.
(age=20
height=NULL
weight=200)!

Rule in SQL: include only tuples that yield TRUE (1.0)

Null Values

Can test for NULL explicitly:

- $x \text{ IS NULL}$
- $x \text{ IS NOT NULL}$

```
SELECT *
FROM Person
WHERE age < 25 OR age >= 25
OR age IS NULL
```

Now it includes all!

INNER JOIN:

Product

name	category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

prodName	store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```
SELECT Product.name, Purchase.store
FROM Product
INNER JOIN Purchase
ON Product.name = Purchase.prodName
```



name	store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Products that never sold (with no Purchase tuple) will be lost!

Outer Joins

- An **outer join** returns tuples from the joined relations that don't have a corresponding tuple in the other relations
 - I.e. If we join relations A and B on $a.X = b.X$, and there is an entry in A with $X=5$, but none in B with $X=5$...
 - A LEFT OUTER JOIN will return a tuple (a, NULL)!
- Left outer joins in SQL:

```
SELECT Product.name, Purchase.store
FROM Product
LEFT OUTER JOIN Purchase ON
Product.name = Purchase.prodName
```

Now we'll get products even if they didn't sell

LEFT OUTER JOIN:

Product

name	category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

prodName	store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```
SELECT Product.name, Purchase.store
FROM Product
LEFT OUTER JOIN Purchase
ON Product.name = Purchase.prodName
```



name	store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

Other Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include the both left and right tuples even if there's no match

A relational database

- A relational database schema is a set of relational schemata, one for each relation
- A relational database instance is a set of relational instances, one for each relation

Two conventions:

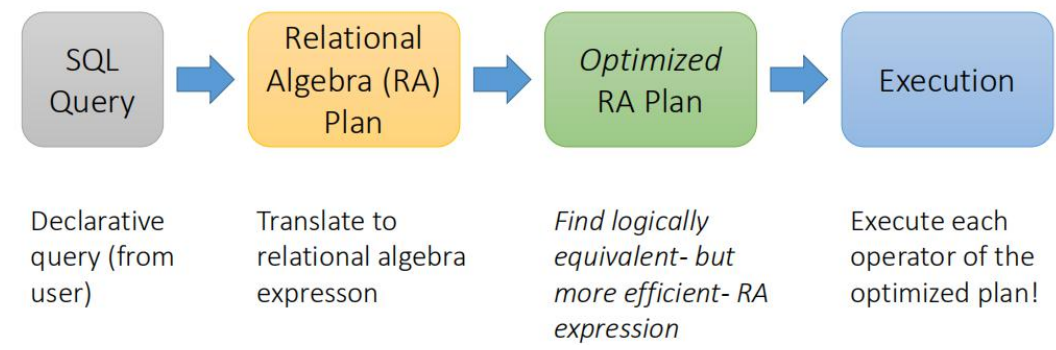
1. We call relational database instances as simply **databases**
2. We assume all instances are valid, i.e., satisfy the domain constraints

Relational Algebra (RA)

- Five basic operators:
 - Selection: σ
 - Projection: Π
 - Cartesian Product: \times
 - Union: \cup
 - Difference: $-$
- Derived or auxiliary operators:
 - Intersection, complement
 - Joins (natural, equi-join, theta join, semi-join)
 - Renaming: ρ
 - Division

RDBMS Architecture

How does an SQL engine work ?



Keep in mind: RA operates on sets!

- RDBMSs use *multisets*, however in relational algebra formalism we will consider sets!
- Also: we will consider the **named perspective**, where every attribute must have a unique name
 - \rightarrow attribute order does not matter...

Selection (σ)

- Returns all tuples which satisfy a condition
- Notation: $\sigma_c(R)$
- Examples
 - $\sigma_{\text{Salary} > 40000}(\text{Employee})$
 - $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$
- The condition c can be
 - $=, <, \leq, >, \geq, <>$

Students(sid,sname,gpa)

SQL:

```
SELECT *  
FROM Students  
WHERE gpa > 3.5;
```



RA:

$\sigma_{gpa > 3.5}(\text{Students})$

Projection (Π)

- Eliminates columns, then removes duplicates
- Notation: $\Pi_{A1, \dots, An}(R)$
- Example: project social-security number and names:
 - $\Pi_{SSN, Name}(\text{Employee})$
 - Output schema: *Answer (SSN, Name)*

Students(sid,sname,gpa)

SQL:

```
SELECT DISTINCT  
sname,  
gpa  
FROM Students;
```



RA:

$\Pi_{sname, gpa}(\text{Students})$

Cross-Product (\times)

- Each tuple in R1 with each tuple in R2
- Notation: $R1 \times R2$
- Example:
 - $\text{Employee} \times \text{Departments}$
- Mainly used to express joins

Students(sid,sname,gpa)
People(ssn,pname,address)

SQL:

```
SELECT *  
FROM Students, People;
```



RA:

$\text{Students} \times \text{People}$

Renaming (ρ – *Rho*)

- Changes the schema, not the instance
- A 'special' operator- neither basic nor derived
- Notation: $\rho_{B1, \dots, Bn}(R)$
- **Note: this is shorthand for the proper form (since names, not order matters!):**
 - $\rho_{A1 \rightarrow B1, \dots, An \rightarrow Bn}(R)$

Students(sid,sname,gpa)

SQL:

```
SELECT  
sid AS studId,  
sname AS name,  
gpa AS gradePtAvg  
FROM Students;
```



RA:

$\rho_{studId, name, gradePtAvg}(\text{Students})$

We care about this operator *because* we are working in a *named perspective*

Natural Join (\bowtie)

- Notation: $R_1 \bowtie R_2$
- Joins R_1 and R_2 on *equality of all shared attributes*
 - If R_1 has attribute set A , and R_2 has attribute set B , and they share attributes $A \cap B = C$, can also be written: $R_1 \bowtie_C R_2$
- Our first example of a *derived* RA operator:
 - Meaning: $R_1 \bowtie R_2 = \Pi_{A \cup B}(\sigma_{C=D}(\rho_{C \rightarrow D}(R_1) \times R_2))$
 - Where:
 - The rename $\rho_{C \rightarrow D}$ renames the shared attributes in one of the relations
 - The selection $\sigma_{C=D}$ checks equality of the shared attributes
 - The projection $\Pi_{A \cup B}$ eliminates the duplicate common attributes

Students(sid,name,gpa)
People(ssn,name,address)

SQL:

```
SELECT DISTINCT
  ssid, S.name, gpa,
  ssn, address
FROM
  Students S,
  People P
WHERE S.name = P.name;
```

RA:

$Students \bowtie People$

Example: Converting SFW Query -> RA

Students(sid,sname,gpa)
People(ssn,sname,address)

```
SELECT DISTINCT
  gpa,
  address
FROM Students S,
  People P
WHERE gpa > 3.5 AND
  S.sname = P.sname;
```



$\Pi_{gpa,address}(\sigma_{gpa>3.5}(S \bowtie P))$

How do we represent
this query in RA?

Division

- Notation: $r \div s$
- It has nothing to do with arithmetic division.
- Let r and s be relations on schemas R and S respectively where

• $R = (A_1, \dots, A_m, B_1, \dots, B_n)$

• $S = (B_1, \dots, B_n)$

The result of $r \div s$ is a relation on schema

$R - S = (A_1, \dots, A_m)$

$r \div s = \{t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r)\}$

Where tu means the concatenation of tuples t and u to produce a single tuple

Division Operation - Example

Relations r, s

A	B	C	D	E
α	a	α	a	1
α	a	γ	a	1
α	a	γ	b	1
β	a	γ	a	1
β	a	γ	b	3
γ	a	γ	a	1
γ	a	γ	b	1
γ	a	β	b	1

r

D	E
a	1
b	1


s

$r \div s$

A	B	C
α	a	γ
γ	a	γ


Union (\cup) and Difference ($-$)

- $R1 \cup R2$
- Example:
 - $\text{ActiveEmployees} \cup \text{RetiredEmployees}$
- $R1 - R2$
- Example:
 - $\text{AllEmployees} - \text{RetiredEmployees}$



What about Intersection (\cap) ?

- It is a derived operator
- $R1 \cap R2 = R1 - (R1 - R2)$
- Also expressed as a join!
- Example
 - $\text{UnionizedEmployees} \cap \text{RetiredEmployees}$



Theta Join (\bowtie_{θ})

- A join that involves a predicate
- $R1 \bowtie_{\theta} R2 = \sigma_{\theta}(R1 \times R2)$
- Here θ can be any condition

```
Students(sid,sname,gpa)
People(ssn,sname,address)
```

SQL:

```
SELECT *
FROM
  Students, People
WHERE  $\theta$ ;
```

Note that natural join is a
theta join + a projection.



RA:
 $\text{Students} \bowtie_{\theta} \text{People}$

Equi-join ($\bowtie_{A=B}$)

- A theta join where θ is an equality
- $R1 \bowtie_{A=B} R2 = \sigma_{A=B}(R1 \times R2)$
- Example:
 - $\text{Employee} \bowtie_{\text{SSN}=\text{SSN}} \text{Dependents}$

```
Students(sid,sname,gpa)
People(ssn,pname,address)
```

SQL:

```
SELECT *
FROM
  Students S,
  People P
WHERE sname = pname;
```



RA:
 $S \bowtie_{\text{sname}=\text{pname}} P$

Most common join
in practice!

Semijoin (\bowtie)

- $R \bowtie S = \Pi_{A_1, \dots, A_n} (R \bowtie S)$
- Where A_1, \dots, A_n are the attributes in R
- Example:
 - $\text{Employee} \bowtie \text{Dependents}$

Students(sid,sname,gpa)
People(ssn,pname,address)

SQL:

```
SELECT DISTINCT
  sid,sname,gpa
FROM
  Students,People
WHERE
  sname = pname;
```


RA:

$\text{Students} \bowtie \text{People}$

Send less data to
reduce network
bandwidth!

Semijoins in Distributed Databases

- Semijoins are often used to compute natural joins in distributed databases



$\text{Employee} \bowtie_{\text{ssn}=\text{ssn}} (\sigma_{\text{age}>71} (\text{Dependents}))$

$R = \text{Employee} \bowtie T$
 $T = \Pi_{\text{ssn}} \sigma_{\text{age}>71} (\text{Dependents})$
 $\text{Answer} = R \bowtie \text{Dependents}$

EXAMPLE RELATIONAL ALGEBRA

Example: Banking Database

branch (branch_name, branch_city, assets)

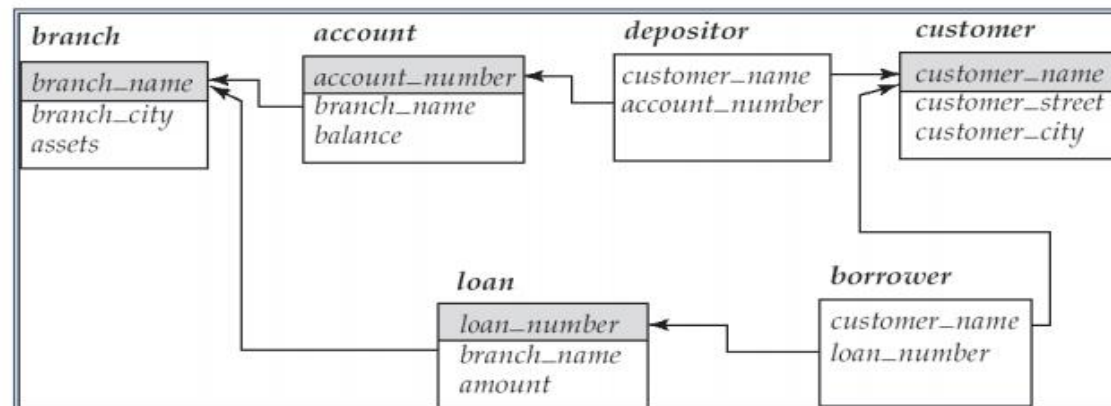
customer (customer_name, customer_street, customer_city)

account (account_number, branch_name, balance)

loan (loan_number, branch_name, amount)

depositor (customer_name, account_number)

borrower (customer_name, loan_number)



- Find all loans of over \$1200

$$\sigma_{amount > 1200} (loan)$$

- Find the loan number for each loan of an amount greater than \$1200

$$\Pi_{loan_number} (\sigma_{amount > 1200} (loan))$$

- Find the names of all customers who have a loan at the Perryridge branch.

$$\Pi_{customer_name} (\sigma_{branch_name = "Perryridge"} (\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan)))$$

- Find the names of all customers who have a loan at the Perryridge branch but do not have an account at any branch of the bank.

$$\Pi_{customer_name} (\sigma_{branch_name = "Perryridge"} (\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan))) - \Pi_{customer_name} (depositor)$$

- Find all customers who have an account at all branches located in Brooklyn city.

$$\Pi_{customer_name, branch_name} (depositor \bowtie account) \div \Pi_{branch_name} (\sigma_{branch_city = "Brooklyn"} (branch))$$

- Find the names of all customers who have a loan, an account, or both, from the bank

$$\Pi_{customer_name} (borrower) \cup \Pi_{customer_name} (depositor)$$

- Find the names of all customers who have a loan and an account at bank.

$$\Pi_{customer_name} (borrower) \cap \Pi_{customer_name} (depositor)$$

- Find the names of all customers who have a loan at the Perryridge branch.

Query 1

$$\Pi_{customer_name} (\sigma_{branch_name = "Perryridge"} (\sigma_{borrower.loan_number = loan.loan_number} (borrower \times loan)))$$

Query 2

$$\Pi_{customer_name} (\sigma_{loan.loan_number = borrower.loan_number} (\sigma_{branch_name = "Perryridge"} (loan) \times borrower))$$